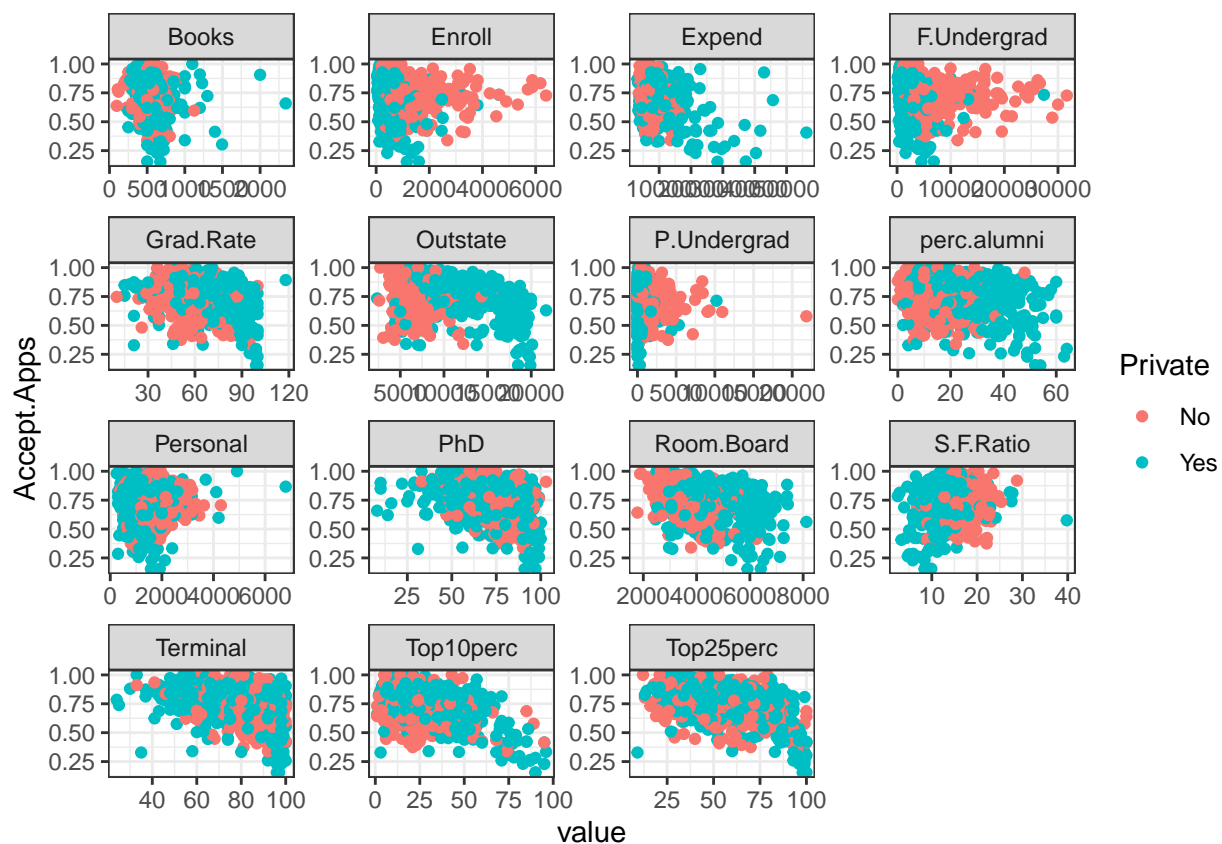


Q3.

(a)

```
library(ISLR)
data("College")
College$Accept.Apps<-College$Accept/College$Apps
College<-College[,c(-2,-3)]

#Scatterplot for variables except for private, private shown as color
College %>%
  gather(-Accept.Apps, -Private, key = "var", value = "value") %>%
  ggplot(aes(x = value, y = Accept.Apps, color = Private)) +
    geom_point() +
    facet_wrap(~ var, scales = "free") +
    theme_bw()
```



```
#Split the dataset
set.seed(234)
inTrain <- createDataPartition(College$Accept.Apps, p = 0.7, list = FALSE)
training <- College[inTrain,]
testing <- College[-inTrain,]
```

(b) Fit a linear model using least squares on the training set, and report the training and test error obtained, with Accept/Apps as the response variable and all the other variables except Accept and Apps as predictors.

For testing error

```
linear_model <- lm(Accept.Apps~. ,data=training)
testPrediction <- predict(linear_model, testing)
test_MSE<-mean((testPrediction - testing$Accept.Apps)^2)
test_MSE
```

```
## [1] 0.01350159
```

```
##Hence, the testing error is 0.01350159
```

For training error

```
trainPrediction <- predict(linear_model, training)
train_MSE<-mean((trainPrediction - training$Accept.Apps)^2)
train_MSE
```

```
## [1] 0.01422197
```

```
##Hence, the testing error is 0.01422197
```

(c)

```
library(leaps)
```

```
## Warning: package 'leaps' was built under R version 4.0.3
```

```
regfit.full = regsubsets(Accept.Apps~. , data = College, nvmax=ncol(College)-1)
regfit.Summary = summary(regfit.full)
names(regfit.Summary)
```

```
## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
```

```
regfit.Summary$rsq
```

```
## [1] 0.2291301 0.2566215 0.2978282 0.3075737 0.3153663 0.3228179 0.3303928
## [8] 0.3371194 0.3451705 0.3525302 0.3542739 0.3549540 0.3556430 0.3563261
## [15] 0.3565993 0.3566201
```

#Adjusted R-square

```
best_adjr2 = which.max(regfit.Summary$adjr2)
best_adjr2 #11
```

```
## [1] 11
```

```
# We can use the coef() function to see which predictors made the cut
coef(regfit.full, 11)
```

```
## (Intercept) PrivateYes Enroll Top10perc P.Undergrad
## 1.072220e+00 7.355125e-02 2.759506e-05 -3.158082e-03 -1.061382e-05
## Outstate Room.Board Books S.F.Ratio perc.alumni
## 5.321032e-06 -2.509935e-05 -8.008527e-05 -4.224509e-03 6.559062e-04
## Expend Grad.Rate
## -7.195160e-06 -1.409409e-03
```

For testing error

```
linear_model1 <- lm(Accept.Apps ~ Private+Enroll+Top10perc+P.Undergrad+Outstate+Room.Board+ Books +S.
testPrediction1 <- predict(linear_model1, testing)
test_MSE1<-mean((testPrediction1 - testing$Accept.Apps)^2)
test_MSE1
```

```
## [1] 0.01327613
```

```
##Hence, the testing error is 0.01327613
```

For training error

```
trainPrediction1 <- predict(linear_model1, training)
train_MSE1<-mean((trainPrediction1 - training$Accept.Apps)^2)
train_MSE1
```

```
## [1] 0.01433002
```

```
##Hence, the testing error is 0.01433002
```

AIC

```
best_cp = which.min(regfit.Summary$cp)
best_cp #11
```

```
## [1] 11
```

```
coef(regfit.full, 11)
```

```
## (Intercept) PrivateYes Enroll Top10perc P.Undergrad
## 1.072220e+00 7.355125e-02 2.759506e-05 -3.158082e-03 -1.061382e-05
## Outstate Room.Board Books S.F.Ratio perc.alumni
## 5.321032e-06 -2.509935e-05 -8.008527e-05 -4.224509e-03 6.559062e-04
## Expend Grad.Rate
## -7.195160e-06 -1.409409e-03
```

For testing error

```
linear_model2 <- lm(Accept.Apps ~ Private+Enroll+Top10perc+P.Undergrad+Outstate+Room.Board+ Books +S.
testPrediction2 <- predict(linear_model2, testing)
test_MSE2<-mean((testPrediction2 - testing$Accept.Apps)^2)
test_MSE2
```

```
## [1] 0.01327613
```

```
##Hence, the testing error is 0.01327613
```

For training error

```
trainPrediction2 <- predict(linear_model2, training)
train_MSE2<-mean((trainPrediction2 - training$Accept.Apps)^2)
train_MSE2
```

```
## [1] 0.01433002
```

```
##Hence, the testing error is 0.01433002
```

BIC

```
best_bic = which.min(regfit.Summary$bic)
best_bic
```

```
## [1] 10
```

```
best_bic #10
```

```
## [1] 10
```

```
coef(regfit.full, 10)
```

```
## (Intercept) PrivateYes      Enroll      Top10perc  P.Undergrad  
## 1.078867e+00 7.523625e-02 2.690366e-05 -3.068681e-03 -1.077537e-05  
##      Outstate      Room.Board      Books      S.F.Ratio      Expend  
## 5.992015e-06 -2.634089e-05 -8.232725e-05 -4.355397e-03 -7.172649e-06  
##      Grad.Rate  
## -1.309782e-03
```

For testing error

```
linear_model3 <- lm(Accept.Apps ~ Private+Enroll+Top10perc+P.Undergrad+Outstate+Room.Board+ Books +S.  
testPrediction3 <- predict(linear_model3, testing)  
test_MSE3<-mean((testPrediction3 - testing$Accept.Apps)^2)  
test_MSE3
```

```
## [1] 0.01315906
```

```
##Hence, the testing error is 0.01315906
```

For training error

```
trainPrediction3 <- predict(linear_model3, training)  
train_MSE3<-mean((trainPrediction3 - training$Accept.Apps)^2)  
train_MSE3
```

```
## [1] 0.01441521
```

```
##Hence, the testing error is 0.01441521
```

(d) #Candidate model in (c) The model with smaller test error is the model founded by BIC with 10 variables included.

```
y_train <- training$Accept.Apps  
y_test <- testing$Accept.Apps  
  
one_hot_encoding <- dummyVars(Accept.Apps ~ Private+Enroll+Top10perc+P.Undergrad+Outstate+Room.Board+ B  
x_train <- predict(one_hot_encoding, training)  
x_test <- predict(one_hot_encoding, testing)
```

```
linear_fit <- train(x = x_train, y = y_train,  
  method = 'lm',  
  trControl = trainControl(method = 'cv', number = 5),  
)
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit  
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit  
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit  
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
(linear_info_test <- postResample(predict(linear_fit, x_test), y_test))
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
##      RMSE    Rsquared      MAE
## 0.11471296 0.30483774 0.09073556
```

```
test_MSE4<-linear_info_test[1]^2
test_MSE4 #0.01315906
```

```
##      RMSE
## 0.01315906
```

```
(linear_info_train <- postResample(predict(linear_fit, x_train), y_train))
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
##      RMSE    Rsquared      MAE
## 0.12006336 0.36615176 0.09392975
```

```
train_MSE4<-linear_info_train[1]^2
train_MSE4 #0.01441521
```

```
##      RMSE
## 0.01441521
```

```
college.glm <- glm(Accept.Apps ~ Private+Enroll+Top10perc+P.Undergrad+Outstate+Room.Board+ Books+S.F.Ra
cv.err = cv.glm(College ,college.glm , K = 5)$delta
cv.err
```

```
## [1] 0.01450695 0.01444841
```

Hence, the test error and train error obtained from cross validation is the same as the value we calculated in (c).

#Full_model in (b)

```
y_train <- training$Accept.Apps
y_test <- testing$Accept.Apps
```

```
one_hot_encoding <- dummyVars(Accept.Apps ~. , data = training)
x_train <- predict(one_hot_encoding, training)
x_test <- predict(one_hot_encoding, testing)
```

```
linear_fit <- train(x = x_train, y = y_train,
                    method = 'lm',
                    trControl = trainControl(method = 'cv', number = 5),
                    )
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
(linear_info_test <- postResample(predict(linear_fit, x_test), y_test))
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
##          RMSE    Rsquared          MAE
## 0.11619634 0.28910308 0.09202104
```

```
test_MSE5<-linear_info_test[1]^2
test_MSE5 #0.01350159
```

```
##          RMSE
## 0.01350159
```

```
(linear_info_train <- postResample(predict(linear_fit, x_train), y_train))
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
##          RMSE    Rsquared          MAE
## 0.11925591 0.37464865 0.09349292
```

```
train_MSE5<-linear_info_train[1]^2
train_MSE5 #0.01422197
```

```
##          RMSE
## 0.01422197
```

```
college.glm <- glm(Accept.Apps ~., data = College)
cv.err = cv.glm(College ,college.glm , K = 5)$delta
cv.err
```

```
## [1] 0.01454782 0.01447501
```

Hence, the test error and train error obtained from cross validation is the same as the value we calculated in (b). (e)

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.0.3
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##      expand, pack, unpack
```

```
## Loaded glmnet 4.0-2

set.seed(1)
grid = 10^seq(10, -2, length=100)
ridge.mod = glmnet(x_train, y_train, alpha=0, lambda=grid)

cv.out = cv.glmnet(x_train, y_train, alpha=0, lambda = grid)

bestlam = cv.out$lambda.min
bestlam

## [1] 0.01

# training MSE
ridge.pred_train = predict(ridge.mod, s=bestlam, newx=x_train)
train_MSE_ridge<-mean((ridge.pred_train-y_train)^2)
train_MSE_ridge

## [1] 0.01432566

# test MSE
ridge.pred_test = predict(ridge.mod, s=bestlam, newx=x_test)
test_MSE_ridge<-mean((ridge.pred_test-y_test)^2)
test_MSE_ridge

## [1] 0.01334943

#cross_validation error
cv.out = cv.glmnet(x_train, y_train, alpha=0, lambda = grid, nfolds = 10)
lambda.grid = cv.out$lambda # grid of lambdas used by cv.glmnet()
mses = cv.out$cvm # mean crossvalidated error (MSE) for each lambda (averaged over the 10 folds)
cv_error = mses[which(lambda.grid == bestlam)] # this is the crossvalidated error (MSE)
print(cv_error)

## [1] 0.01538698

(f)
lasso.mod = glmnet(x_train, y_train, alpha=1, lambda=grid)
set.seed(1)
cv.out = cv.glmnet(x_train, y_train, alpha=1)
# best lambda
bestlam = cv.out$lambda.min
bestlam

## [1] 0.0009319414

# training error
lasso.pred_train = predict(lasso.mod,s=bestlam,newx=x_train)
train_MSE_lasso<-mean((lasso.pred_train-y_train)^2)
train_MSE_lasso

## [1] 0.01554418

# test error
lasso.pred_test = predict(lasso.mod,s=bestlam,newx=x_test)
test_MSE_lasso<-mean((lasso.pred_test-y_test)^2)
test_MSE_lasso

## [1] 0.01405796
```

```

cv.out = cv.glmnet(x_train, y_train, alpha=1, lambda = grid, nfolds = 10)
lambda.grid = cv.out$lambda =
msep = cv.out$cvm =
cv_error = msep[which(lambda.grid == bestlam)]
print(cv_error)

```

```
## numeric(0)
```

```
(g)
```

```

train_MSE_best_reduced<-train_MSE3
test_MSE_best_reduced<-test_MSE3

```

```

models = c("Best reduced OLS", "Ridge Regression", "Lasso")
train_err = c(
  train_MSE_best_reduced,
  train_MSE_ridge,
  train_MSE_lasso)

```

```

test_err = c(
  test_MSE_best_reduced,
  test_MSE_ridge,
  test_MSE_lasso
)

```

```

results = data.frame(
  models,
  train_err,
  test_err
)
colnames(results) = c("Model", "Train MSE", "Test MSE")
print(results)

```

```

##           Model  Train MSE   Test MSE
## 1 Best reduced OLS 0.01441521 0.01315906
## 2 Ridge Regression 0.01432566 0.01334943
## 3           Lasso 0.01554418 0.01405796

```

In all, we can predict the acceptance rate with approximately 0.0135 testing MSE. The testing MSE with ridge regression is the lowest, best reduced OLS model is the second lowest and the Lasso regression is the highest. In this dataset, I would choose ridge regression, since it reports the lowest training and test error.