

Q2.

```
library(e1071)

## Warning: package 'e1071' was built under R version 4.0.3

data(crabs)
head(crabs)

##   sp sex index  FL  RW  CL  CW  BD
## 1  B  M     1  8.1 6.7 16.1 19.0 7.0
## 2  B  M     2  8.8 7.7 18.1 20.8 7.4
## 3  B  M     3  9.2 7.8 19.0 22.4 7.7
## 4  B  M     4  9.6 7.9 20.1 23.1 8.2
## 5  B  M     5  9.8 8.0 20.3 23.0 8.2
## 6  B  M     6 10.8 9.0 23.0 26.5 9.8

set.seed(6789)
inTrain <- createDataPartition(crabs$sp, p = 0.8, list = FALSE)
training <- crabs[inTrain,]
testing <- crabs[-inTrain,]
```

a

For the linear Svm model

```
set.seed(1)

ranges = c(0.001, 0.01, 0.1, 1, 5, 10, 100)

svmTrainErr = vector(length = length(ranges))
svmTestErr = vector(length = length(ranges))

for (i in 1:length(ranges)) {
  svm.model = svm(sp ~ .-index, data = training, kernel = "linear", cost = ranges[i], scale = FALSE)
  pred.train <- predict(svm.model, training)
  pred.test <- predict(svm.model, testing)
  svmTrainErr[i] = mean(pred.train != training$sp)
  svmTestErr[i] = mean(pred.test != testing$sp)
}

tune.out <- tune(svm, sp ~ .-index, data = training, kernel = "linear",
  ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0
```

```
##
## - Detailed performance results:
##   cost   error dispersion
## 1 1e-03 0.56250 0.12842529
## 2 1e-02 0.35625 0.11043632
## 3 1e-01 0.13750 0.07095578
## 4 1e+00 0.00000 0.00000000
## 5 5e+00 0.00000 0.00000000
## 6 1e+01 0.00000 0.00000000
## 7 1e+02 0.00000 0.00000000
```

```
tune.out$best.model
```

```
##
## Call:
## best.tune(method = svm, train.x = sp ~ . - index, data = training,
##   ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##         cost: 1
##
## Number of Support Vectors: 49
```

```
svmTestErr
```

```
## [1] 0.4 0.0 0.0 0.0 0.0 0.0 0.0
```


```
svmTrainErr
```

```
## [1] 0.31875 0.00625 0.00000 0.00000 0.00000 0.00000 0.00000
```

We take the logarithm of the cost to make plots for better visualization


```
plot(log(tune.out$performances$cost), tune.out$performances$error, main = "Cross-validation error vs. cost")
```

Cross-validation error vs. cost




```
plot(log(ranges),svmTrainErr,main = "Train error vs. cost", xlab = "log(cost)",ylab = "Error" )
```

Train error vs. cost



```
plot(log(ranges),svmTestErr,main = "Test error vs. cost", xlab = "log(cost)",ylab = "Error" )
```



The best model is obtained by $cost \geq 1$, Since the cv-error, train error and test error are all minimized.

b

For the non-linear Svm model

```
ranges = c(0.1, 1, 10, 100, 1000)

set.seed(1)
tune.out <- tune(svm,sp ~ .-index, data = training, kernel = "radial",
ranges = list(cost = c(0.1, 1, 10, 100, 1000),
gamma = c(0.5, 1, 2, 3, 4),degree = c(1,2,3,4,5)))
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma degree
##   10   0.5      1
##
## - best performance: 0.00625
##
## - Detailed performance results:
##   cost gamma degree  error dispersion
## 1  1e-01  0.5      1 0.33750 0.15080801
```


## 2	1e+00	0.5	1	0.07500	0.06454972
## 3	1e+01	0.5	1	0.00625	0.01976424
## 4	1e+02	0.5	1	0.00625	0.01976424
## 5	1e+03	0.5	1	0.00625	0.01976424
## 6	1e-01	1.0	1	0.43750	0.13501543
## 7	1e+00	1.0	1	0.06250	0.05892557
## 8	1e+01	1.0	1	0.01250	0.02635231
## 9	1e+02	1.0	1	0.01250	0.02635231
## 10	1e+03	1.0	1	0.01250	0.02635231
## 11	1e-01	2.0	1	0.53750	0.17969882
## 12	1e+00	2.0	1	0.06875	0.05472469
## 13	1e+01	2.0	1	0.02500	0.03227486
## 14	1e+02	2.0	1	0.02500	0.03227486
## 15	1e+03	2.0	1	0.02500	0.03227486
## 16	1e-01	3.0	1	0.53125	0.20252315
## 17	1e+00	3.0	1	0.06250	0.05892557
## 18	1e+01	3.0	1	0.03125	0.03294039
## 19	1e+02	3.0	1	0.03125	0.03294039
## 20	1e+03	3.0	1	0.03125	0.03294039
## 21	1e-01	4.0	1	0.51875	0.22831280
## 22	1e+00	4.0	1	0.05625	0.06215181
## 23	1e+01	4.0	1	0.03125	0.03294039
## 24	1e+02	4.0	1	0.03125	0.03294039
## 25	1e+03	4.0	1	0.03125	0.03294039
## 26	1e-01	0.5	2	0.33750	0.15080801
## 27	1e+00	0.5	2	0.07500	0.06454972
## 28	1e+01	0.5	2	0.00625	0.01976424
## 29	1e+02	0.5	2	0.00625	0.01976424
## 30	1e+03	0.5	2	0.00625	0.01976424
## 31	1e-01	1.0	2	0.43750	0.13501543
## 32	1e+00	1.0	2	0.06250	0.05892557
## 33	1e+01	1.0	2	0.01250	0.02635231
## 34	1e+02	1.0	2	0.01250	0.02635231
## 35	1e+03	1.0	2	0.01250	0.02635231
## 36	1e-01	2.0	2	0.53750	0.17969882
## 37	1e+00	2.0	2	0.06875	0.05472469
## 38	1e+01	2.0	2	0.02500	0.03227486
## 39	1e+02	2.0	2	0.02500	0.03227486
## 40	1e+03	2.0	2	0.02500	0.03227486
## 41	1e-01	3.0	2	0.53125	0.20252315
## 42	1e+00	3.0	2	0.06250	0.05892557
## 43	1e+01	3.0	2	0.03125	0.03294039
## 44	1e+02	3.0	2	0.03125	0.03294039
## 45	1e+03	3.0	2	0.03125	0.03294039
## 46	1e-01	4.0	2	0.51875	0.22831280
## 47	1e+00	4.0	2	0.05625	0.06215181
## 48	1e+01	4.0	2	0.03125	0.03294039
## 49	1e+02	4.0	2	0.03125	0.03294039
## 50	1e+03	4.0	2	0.03125	0.03294039
## 51	1e-01	0.5	3	0.33750	0.15080801
## 52	1e+00	0.5	3	0.07500	0.06454972
## 53	1e+01	0.5	3	0.00625	0.01976424
## 54	1e+02	0.5	3	0.00625	0.01976424
## 55	1e+03	0.5	3	0.00625	0.01976424

## 56	1e-01	1.0	3	0.43750	0.13501543
## 57	1e+00	1.0	3	0.06250	0.05892557
## 58	1e+01	1.0	3	0.01250	0.02635231
## 59	1e+02	1.0	3	0.01250	0.02635231
## 60	1e+03	1.0	3	0.01250	0.02635231
## 61	1e-01	2.0	3	0.53750	0.17969882
## 62	1e+00	2.0	3	0.06875	0.05472469
## 63	1e+01	2.0	3	0.02500	0.03227486
## 64	1e+02	2.0	3	0.02500	0.03227486
## 65	1e+03	2.0	3	0.02500	0.03227486
## 66	1e-01	3.0	3	0.53125	0.20252315
## 67	1e+00	3.0	3	0.06250	0.05892557
## 68	1e+01	3.0	3	0.03125	0.03294039
## 69	1e+02	3.0	3	0.03125	0.03294039
## 70	1e+03	3.0	3	0.03125	0.03294039
## 71	1e-01	4.0	3	0.51875	0.22831280
## 72	1e+00	4.0	3	0.05625	0.06215181
## 73	1e+01	4.0	3	0.03125	0.03294039
## 74	1e+02	4.0	3	0.03125	0.03294039
## 75	1e+03	4.0	3	0.03125	0.03294039
## 76	1e-01	0.5	4	0.33750	0.15080801
## 77	1e+00	0.5	4	0.07500	0.06454972
## 78	1e+01	0.5	4	0.00625	0.01976424
## 79	1e+02	0.5	4	0.00625	0.01976424
## 80	1e+03	0.5	4	0.00625	0.01976424
## 81	1e-01	1.0	4	0.43750	0.13501543
## 82	1e+00	1.0	4	0.06250	0.05892557
## 83	1e+01	1.0	4	0.01250	0.02635231
## 84	1e+02	1.0	4	0.01250	0.02635231
## 85	1e+03	1.0	4	0.01250	0.02635231
## 86	1e-01	2.0	4	0.53750	0.17969882
## 87	1e+00	2.0	4	0.06875	0.05472469
## 88	1e+01	2.0	4	0.02500	0.03227486
## 89	1e+02	2.0	4	0.02500	0.03227486
## 90	1e+03	2.0	4	0.02500	0.03227486
## 91	1e-01	3.0	4	0.53125	0.20252315
## 92	1e+00	3.0	4	0.06250	0.05892557
## 93	1e+01	3.0	4	0.03125	0.03294039
## 94	1e+02	3.0	4	0.03125	0.03294039
## 95	1e+03	3.0	4	0.03125	0.03294039
## 96	1e-01	4.0	4	0.51875	0.22831280
## 97	1e+00	4.0	4	0.05625	0.06215181
## 98	1e+01	4.0	4	0.03125	0.03294039
## 99	1e+02	4.0	4	0.03125	0.03294039
## 100	1e+03	4.0	4	0.03125	0.03294039
## 101	1e-01	0.5	5	0.33750	0.15080801
## 102	1e+00	0.5	5	0.07500	0.06454972
## 103	1e+01	0.5	5	0.00625	0.01976424
## 104	1e+02	0.5	5	0.00625	0.01976424
## 105	1e+03	0.5	5	0.00625	0.01976424
## 106	1e-01	1.0	5	0.43750	0.13501543
## 107	1e+00	1.0	5	0.06250	0.05892557
## 108	1e+01	1.0	5	0.01250	0.02635231
## 109	1e+02	1.0	5	0.01250	0.02635231

```
## 110 1e+03 1.0      5 0.01250 0.02635231
## 111 1e-01 2.0      5 0.53750 0.17969882
## 112 1e+00 2.0      5 0.06875 0.05472469
## 113 1e+01 2.0      5 0.02500 0.03227486
## 114 1e+02 2.0      5 0.02500 0.03227486
## 115 1e+03 2.0      5 0.02500 0.03227486
## 116 1e-01 3.0      5 0.53125 0.20252315
## 117 1e+00 3.0      5 0.06250 0.05892557
## 118 1e+01 3.0      5 0.03125 0.03294039
## 119 1e+02 3.0      5 0.03125 0.03294039
## 120 1e+03 3.0      5 0.03125 0.03294039
## 121 1e-01 4.0      5 0.51875 0.22831280
## 122 1e+00 4.0      5 0.05625 0.06215181
## 123 1e+01 4.0      5 0.03125 0.03294039
## 124 1e+02 4.0      5 0.03125 0.03294039
## 125 1e+03 4.0      5 0.03125 0.03294039
```


We first pick the value of gamma

```
with(tune.out$performances, {
plot(error[gamma == 0.5] ~ cost[gamma == 0.5], ylim = c(.1, .35),
type = "o", col = rainbow(5)[1], ylab = "CV error", xlab = "cost")
lines(error[gamma == 1] ~ cost[gamma == 1],
type = "o", col = rainbow(5)[2])
lines(error[gamma == 2] ~ cost[gamma == 2],
type = "o", col = rainbow(5)[3])
lines(error[gamma == 3] ~ cost[gamma == 3],
type = "o", col = rainbow(5)[4])
lines(error[gamma == 4] ~ cost[gamma == 4],
type = "o", col = rainbow(5)[5])
})
legend("bottom", horiz = T, legend = c(0.5, 1:4), col = rainbow(5),
lty = 1, cex = .75, title = "gamma")
```

Hence the best choice of gamma is 0.5. Then we pick the value of degree

```
with(tune.out$performances, {
  plot(error[degree == 5] ~ cost[degree == 5], ylim = c(.1, .35),
  type = "o", col = rainbow(5)[1], ylab = "CV error", xlab = "cost")
  lines(error[degree == 1] ~ cost[degree == 1],
  type = "o", col = rainbow(5)[2])
  lines(error[degree == 2] ~ cost[degree == 2],
  type = "o", col = rainbow(5)[3])
  lines(error[degree == 3] ~ cost[degree == 3],
  type = "o", col = rainbow(5)[4])
  lines(error[degree == 4] ~ cost[degree == 4],
  type = "o", col = rainbow(5)[5])
})
legend("bottom", horiz = T, legend = c(1:5), col = rainbow(5),
lty = 1, cex = .75, title = "degree")
```



Hence the best choice of degree is 1.

```
bestmod <- tune.out$best.model
summary(bestmod)
```

```
##
## Call:
## best.tune(method = svm, train.x = sp ~ . - index, data = training,
##   ranges = list(cost = c(0.1, 1, 10, 100, 1000), gamma = c(0.5,
##     1, 2, 3, 4), degree = c(1, 2, 3, 4, 5)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:  10
##
## Number of Support Vectors:  43
##
## ( 21 22 )
##
##
## Number of Classes:  2
##
## Levels:
## B 0
```

We take the lograithm of the cost to make plots for better visualization

```

ranges = c(0.1, 1, 10, 100, 1000)


svmTrainErr = vector(length = length(ranges))
svmTestErr = vector(length = length(ranges))

for (i in 1:length(ranges)) {
  svm.model = svm(sp ~ .-index, data = training, kernel = "radial", cost = ranges[i], degree = 1, gamma = 0.001)
  pred.train <- predict(svm.model, training)
  pred.test <- predict(svm.model, testing)
  svmTrainErr[i] = mean(pred.train != training$sp)
  svmTestErr[i] = mean(pred.test != testing$sp)
}

selected <- tune.out$performances$gamma == 0.5 & tune.out$performances$degree == 1
plot(log(ranges), subset(tune.out$performances, selected)$error, main = "Cross-validation error vs. cost",

```

Cross-validation error vs. cost




```

plot(log(ranges), svmTrainErr, main = "Train error vs. cost", xlab = "log(cost)", ylab = "Error" )


```

Train error vs. cost



```
plot(log(ranges),svmTestErr,main = "Test error vs. cost", xlab = "log(cost)",ylab = "Error" )
```

Test error vs. cost



The best model is obtained by $cost \geq 10$, Since the cv-error, train error and test error are all minimized.

The best model is obtained by $cost \geq 1$, Since the cv-error, train error and test error are all minimized.